

homework2

2025-04-26

```
#Imported all the necessary libraries
```

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
```

```
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
```

```
## ✓ forcats   1.0.0      ✓ stringr    1.5.1
```

```
## ✓ ggplot2    3.5.2      ✓ tibble     3.2.1
```

```
## ✓ lubridate  1.9.4      ✓ tidyr      1.3.1
```

```
## ✓ purrr      1.0.4
```

```
## — Conflicts ————— tidyverse_conflicts() —
```

```
## ✗ dplyr::filter() masks stats::filter()
```

```
## ✗ dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to be  
come errors
```

```
library(e1071)
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
##
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      lift
```

```
library(ggplot2)
```

```
library(dplyr)
```

```
# I've set the working directory and loaded 'nhis_2022.csv' into data
```

```
data <- read.csv("C:/Users/alekh/Downloads/nhis_2022.csv")
```

```
# Exploring the data
```

```
head(data)
```

##	YEAR	SERIAL	STRATA	PSU	NHISHID	REGION	PERNUM	NHISPID	HHX	
## 1	2022	1	143	16	0002022H000001	4	1	0002022H00000110	H000001	
## 2	2022	2	106	53	0002022H000003	3	1	0002022H00000310	H000003	
## 3	2022	2	106	53	0002022H000003	3	2	0002022H00000320	H000003	
## 4	2022	3	134	13	0002022H000006	2	1	0002022H00000610	H000006	
## 5	2022	4	106	53	0002022H000007	3	1	0002022H00000710	H000007	
## 6	2022	4	106	53	0002022H000007	3	2	0002022H00000720	H000007	
##	SAMPWEIGHT	ASTATFLG	CSTATFLG	AGE	SEX	MARSTCUR	EDUC	HOURSWRK	POVERTY	HEIGHT
## 1	8018	1	0	61	1	1	201	45	34	69
## 2	10117	1	0	43	1	1	301	45	37	70
## 3	7933	0	1	12	2	0	0	0	37	60
## 4	2681	1	0	68	1	5	505	0	31	75
## 5	10233	1	0	73	1	1	201	0	32	71
## 6	7712	0	1	16	2	0	0	0	32	65
##	WEIGHT	BMICALC	HINOTCOVE	CANCEREV	CHEARTDIEV	DIABETICEV	HEARTATTEV	STROKEV		
## 1	260	38.4	1	1	1	1	1	1		
## 2	190	27.3	1	1	1	1	1	1		
## 3	96	18.7	1	0	0	1	0	0		
## 4	200	25.0	1	1	1	1	1	1		
## 5	172	24.0	1	1	1	1	1	1		
## 6	106	17.6	1	0	0	1	0	0		
##	ALCANYNO	ALCDAYS	SYR	CIGDAYMO	MOD10DMIN	VIG10DMIN	FRUTNO	VEGENO	JUICEMNO	
## 1	2	104	96	0	0	5	15	0		
## 2	1	52	96	20	0	1	1	1		
## 3	996	996	96	0	0	996	996	996		
## 4	7	364	96	60	0	3	1	0		
## 5	0	0	96	690	0	2	4	0		
## 6	996	996	96	0	0	996	996	996		
##	SALADSNO	BEANNO	SALSAMNO	TOMSAUCEMNO	SODAPNO	FRIESPNNO	SPORDRMNO	FRTDRINKMNO		
## 1	10	5	5	2	0	110	3	0		
## 2	1	1	1	1	0	1	0	0		
## 3	996	996	996	996	996	996	996	996		
## 4	1	1	2	1	1	1	0	2		
## 5	4	2	0	3	30	5	1	0		
## 6	996	996	996	996	996	996	996	996		
##	COFETEAMNO	POTATONO	PIZZANO	HRSLEEP	CVDSHT					
## 1	0	3	2	8	1					
## 2	1	1	1	6	2					
## 3	996	996	996	0	2					
## 4	0	1	1	6	2					
## 5	30	6	2	8	2					
## 6	996	996	996	0	1					

```
summary(data)
```

##	YEAR	SERIAL	STRATA	PSU	
##	Min. :2022	Min. : 1	Min. :100.0	Min. : 1.00	
##	1st Qu.:2022	1st Qu.: 7184	1st Qu.:112.0	1st Qu.: 8.00	
##	Median :2022	Median :14403	Median :126.0	Median : 23.00	
##	Mean :2022	Mean :14419	Mean :125.8	Mean : 30.94	
##	3rd Qu.:2022	3rd Qu.:21648	3rd Qu.:140.0	3rd Qu.: 48.00	
##	Max. :2022	Max. :28854	Max. :151.0	Max. :153.00	
##	NHISHID	REGION	PERNUM	NHISPID	
##	Length:35115	Min. :1.000	Min. :1.000	Length:35115	
##	Class :character	1st Qu.:2.000	1st Qu.:1.000	Class :character	
##	Mode :character	Median :3.000	Median :1.000	Mode :character	
##		Mean :2.712	Mean :1.178		
##		3rd Qu.:4.000	3rd Qu.:1.000		
##		Max. :4.000	Max. :2.000		
##	HHX	SAMPWEIGHT	ASTATFLG	CSTATFLG	
##	Length:35115	Min. : 740	Min. :0.0000	Min. :0.0000	
##	Class :character	1st Qu.: 5095	1st Qu.:1.0000	1st Qu.:0.0000	
##	Mode :character	Median : 7947	Median :1.0000	Median :0.0000	
##		Mean : 9343	Mean :0.7874	Mean :0.2126	
##		3rd Qu.:11777	3rd Qu.:1.0000	3rd Qu.:0.0000	
##		Max. :43112	Max. :1.0000	Max. :1.0000	
##	AGE	SEX	MARSTCUR	EDUC	
##	Min. : 0.0	Min. :1.000	Min. :0.000	Min. : 0.0	
##	1st Qu.: 23.0	1st Qu.:1.000	1st Qu.:1.000	1st Qu.:103.0	
##	Median : 45.0	Median :2.000	Median :1.000	Median :301.0	
##	Mean : 45.3	Mean :1.532	Mean :3.349	Mean :248.4	
##	3rd Qu.: 65.0	3rd Qu.:2.000	3rd Qu.:6.000	3rd Qu.:400.0	
##	Max. :999.0	Max. :9.000	Max. :9.000	Max. :999.0	
##	HOURSWRK	POVERTY	HEIGHT	WEIGHT	
##	Min. : 0.00	Min. :11.00	Min. : 0.00	Min. : 0.0	
##	1st Qu.: 0.00	1st Qu.:24.00	1st Qu.:62.00	1st Qu.:130.0	
##	Median : 0.00	Median :33.00	Median :66.00	Median :165.0	
##	Mean :17.64	Mean :30.28	Mean :60.72	Mean :215.2	
##	3rd Qu.:40.00	3rd Qu.:37.00	3rd Qu.:70.00	3rd Qu.:203.0	
##	Max. :99.00	Max. :37.00	Max. :99.00	Max. :999.0	
##	BMICALC	HINOTCOVE	CANCEREV	CHEARTDIEV	
##	Min. : 11.5	Min. :1.000	Min. :0.000	Min. :0.0000	
##	1st Qu.: 24.0	1st Qu.:1.000	1st Qu.:1.000	1st Qu.:1.0000	
##	Median : 28.3	Median :1.000	Median :1.000	Median :1.0000	
##	Mean :215.8	Mean :1.093	Mean :0.892	Mean :0.8533	
##	3rd Qu.: 36.8	3rd Qu.:1.000	3rd Qu.:1.000	3rd Qu.:1.0000	
##	Max. :996.0	Max. :9.000	Max. :9.000	Max. :9.0000	
##	DIABETICEV	HEARTATTEV	STROKEV	ALCANYNO	ALCDAYSyr
##	Min. :1.000	Min. :0.0000	Min. :0.000	Min. : 0.0	Min. : 0
##	1st Qu.:1.000	1st Qu.:1.0000	1st Qu.:1.000	1st Qu.: 1.0	1st Qu.: 6
##	Median :1.000	Median :1.0000	Median :1.000	Median : 3.0	Median :104
##	Mean :1.092	Mean :0.8247	Mean :0.824	Mean :330.5	Mean :372
##	3rd Qu.:1.000	3rd Qu.:1.0000	3rd Qu.:1.000	3rd Qu.:996.0	3rd Qu.:996
##	Max. :9.000	Max. :9.0000	Max. :9.000	Max. :999.0	Max. :999
##	CIGDAYMO	MOD10DMIN	VIG10DMIN	FRUTNO	
##	Min. : 0.00	Min. : 0.0	Min. : 0.00	Min. : 0.0	
##	1st Qu.:96.00	1st Qu.: 0.0	1st Qu.: 0.00	1st Qu.: 1.0	
##	Median :96.00	Median :20.0	Median : 0.00	Median : 3.0	
##	Mean :94.36	Mean : 34.6	Mean : 16.59	Mean :245.5	
##	3rd Qu.:96.00	3rd Qu.: 45.0	3rd Qu.: 15.00	3rd Qu.: 30.0	

##	Max.	:99.00	Max.	:999.0	Max.	:999.00	Max.	:999.0
##	VEGENO		JUICEMNO		SALADSNO		BEANNO	
##	Min.	: 0.0	Min.	: 0.0	Min.	: 0.0	Min.	: 0.0
##	1st Qu.:	1.0	1st Qu.:	0.0	1st Qu.:	1.0	1st Qu.:	1.0
##	Median :	3.0	Median :	1.0	Median :	3.0	Median :	2.0
##	Mean	:247.2	Mean	:244.1	Mean	:245.1	Mean	:245.6
##	3rd Qu.:	30.0	3rd Qu.:	30.0	3rd Qu.:	30.0	3rd Qu.:	20.0
##	Max.	:999.0	Max.	:999.0	Max.	:999.0	Max.	:999.0
##	SALSAMNO		TOMSAUCEMNO		SODAPNO		FRIESPN0	
##	Min.	: 0.0	Min.	: 0.0	Min.	: 0.0	Min.	: 0.0
##	1st Qu.:	0.0	1st Qu.:	1.0	1st Qu.:	0.0	1st Qu.:	1.0
##	Median :	2.0	Median :	2.0	Median :	1.0	Median :	2.0
##	Mean	:245.3	Mean	:246.5	Mean	:243.1	Mean	:244.6
##	3rd Qu.:	20.0	3rd Qu.:	15.0	3rd Qu.:	30.0	3rd Qu.:	20.0
##	Max.	:999.0	Max.	:999.0	Max.	:999.0	Max.	:999.0
##	SPORDRMNO		FRTDRINKMNO		COFETeamNO		POTATONO	
##	Min.	: 0.0	Min.	: 0.0	Min.	: 0.0	Min.	: 0.0
##	1st Qu.:	0.0	1st Qu.:	0.0	1st Qu.:	0.0	1st Qu.:	1.0
##	Median :	0.0	Median :	0.0	Median :	1.0	Median :	2.0
##	Mean	:242.3	Mean	:242.5	Mean	:243.2	Mean	:245.7
##	3rd Qu.:	15.0	3rd Qu.:	10.0	3rd Qu.:	30.0	3rd Qu.:	20.0
##	Max.	:999.0	Max.	:999.0	Max.	:999.0	Max.	:999.0
##	PIZZANO		HRSLEEP		CVDSHT			
##	Min.	: 0.0	Min.	: 0.000	Min.	:0.000		
##	1st Qu.:	1.0	1st Qu.:	5.000	1st Qu.:	1.000		
##	Median :	2.0	Median :	7.000	Median :	2.000		
##	Mean	:244.8	Mean	: 8.135	Mean	:1.791		
##	3rd Qu.:	10.0	3rd Qu.:	8.000	3rd Qu.:	2.000		
##	Max.	:999.0	Max.	:99.000	Max.	:9.000		

```
str(data)
```

```

## 'data.frame':    35115 obs. of  48 variables:
## $ YEAR          : int  2022 2022 2022 2022 2022 2022 2022 2022 2022 2022 ...
## $ SERIAL        : int   1 2 2 3 4 4 5 6 7 8 ...
## $ STRATA        : int  143 106 106 134 106 106 127 111 143 105 ...
## $ PSU           : int   16 53 53 13 53 53 26 11 14 61 ...
## $ NHISHID       : chr   "0002022H000001" "0002022H000003" "0002022H000003" "0002022H000006"
...
## $ REGION        : int   4 3 3 2 3 3 2 4 4 1 ...
## $ PERNUM        : int   1 1 2 1 1 2 1 1 1 1 ...
## $ NHISPID       : chr   "0002022H00000110" "0002022H00000310" "0002022H00000320" "0002022H000
00610" ...
## $ HHX           : chr   "H000001" "H000003" "H000003" "H000006" ...
## $ SAMPWEIGHT    : num  8018 10117 7933 2681 10233 ...
## $ ASTATFLG      : int   1 1 0 1 1 0 1 1 1 1 ...
## $ CSTATFLG      : int   0 0 1 0 0 1 0 0 0 0 ...
## $ AGE           : int   61 43 12 68 73 16 73 21 59 67 ...
## $ SEX           : int   1 1 2 1 1 2 1 1 1 2 ...
## $ MARSTCUR      : int   1 1 0 5 1 0 1 8 7 7 ...
## $ EDUC          : int  201 301 0 505 201 0 201 303 201 400 ...
## $ HOURSWRK      : int   45 45 0 0 0 0 0 0 0 6 ...
## $ POVERTY       : int   34 37 37 31 32 32 36 23 33 37 ...
## $ HEIGHT        : int   69 70 60 75 71 65 71 68 68 63 ...
## $ WEIGHT        : int  260 190 96 200 172 106 190 200 175 169 ...
## $ BMICALC       : num  38.4 27.3 18.7 25 24 17.6 26.5 30.4 26.6 29.9 ...
## $ HINOTCOVE     : int   1 1 1 1 1 1 1 9 2 1 ...
## $ CANCEREV      : int   1 1 0 1 1 0 1 1 2 2 ...
## $ CHEARTDIEV    : int   1 1 0 1 1 0 2 1 1 2 ...
## $ DIABETICEV    : int   1 1 1 1 1 1 1 1 1 1 ...
## $ HEARTATTEV    : int   1 1 0 1 1 0 1 1 1 2 ...
## $ STROKEV       : int   1 1 0 1 1 0 1 1 1 1 ...
## $ ALCANYNO      : int   2 1 996 7 0 996 2 996 4 997 ...
## $ ALCDAYSyr     : int  104 52 996 364 0 996 2 996 4 997 ...
## $ CIGDAYMO      : int   96 96 96 96 96 96 96 96 96 30 ...
## $ MOD10DMIN     : int   0 20 0 60 690 0 60 45 15 120 ...
## $ VIG10DMIN     : int   0 0 0 0 0 0 0 45 0 0 ...
## $ FRUTNO        : int   5 1 996 3 2 996 1 0 1 1 ...
## $ VEGENO        : int   15 1 996 1 4 996 2 2 0 1 ...
## $ JUICEMNO      : int   0 1 996 0 0 996 10 3 0 0 ...
## $ SALADSNO      : int   10 1 996 1 4 996 5 2 1 3 ...
## $ BEANNO        : int   5 1 996 1 2 996 0 2 0 3 ...
## $ SALSAMNO      : int   5 1 996 2 0 996 0 1 2 1 ...
## $ TOMSAUCEMNO   : int   2 1 996 1 3 996 4 0 1 3 ...
## $ SODAPNO       : int   0 0 996 1 30 996 5 2 0 0 ...
## $ FRIESPN0      : int  110 1 996 1 5 996 3 0 4 0 ...
## $ SPORDRMNO     : int   3 0 996 0 1 996 3 0 1 5 ...
## $ FRTDRINKMNO   : int   0 0 996 2 0 996 3 0 0 0 ...
## $ COFETEAMNO    : int   0 1 996 0 30 996 0 0 0 1 ...
## $ POTATONO      : int   3 1 996 1 6 996 1 0 3 2 ...
## $ PIZZANO       : int   2 1 996 1 2 996 1 1 1 3 ...
## $ HRSLEEP       : int   8 6 0 6 8 0 6 9 9 8 ...
## $ CVDSHT        : int   1 2 2 2 2 1 2 2 1 2 ...

```

```
# Subsetting the data, taking only adults, i.e., between the ages 18 and 70.
# Converting numeric variables to categoric factors.
```

```
data <- data %>%
  filter(AGE >= 18, AGE <= 70, STROKEV %in% c(1,2)) %>%
  mutate(
    Sex = factor(SEX, levels = c(1,2), labels = c("Male","Female")),
    Stroke = factor(STROKEV, levels = c(1,2), labels = c("No","Yes"))
  )
```

```
# Renaming variables to clear column names.
```

```
names(data)[names(data)=="AGE"] <- "Age"
names(data)[names(data)=="HRSLEEP"] <- "Hours Of Sleep"
names(data)[names(data)=="HOURSWRK"] <- "Hours Worked"
names(data)[names(data)=="ALCDAYSyr"] <- "Alcohol Consumption Days Per Year"
names(data)[names(data)=="CIGDAYMO"] <- "Cigarettes Consumed Per Month"
names(data)[names(data)=="MOD10DMIN"] <- "Duration Of Moderate Activity(in mins)"
names(data)[names(data)=="VIG10DMIN"] <- "Duration Of Vigorous Activity(in mins)"
```

```
# Predicting the stroke status (Yes/No) in adults (18-70)
```

```
# using SVMs on predictors:
```

```
# Age, Sex, Hours Of Sleep, Hours Worked,Alcohol Consumption Days
# Cigarettes/Month, Moderate & Vigorous Activity (mins).
```

```
# Cleaning the invalid codes and then replace those with NA,
#lastly, drop those null values.
```


```
codes <- c(996, 997, 998, 999)
variables <- c("Age", "Hours Of Sleep", "Hours Worked",
              "Alcohol Consumption Days Per Year",
              "Cigarettes Consumed Per Month",
              "Duration Of Moderate Activity(in mins)",
              "Duration Of Vigorous Activity(in mins)")
# Keeping both Moderate and Vigorous activity
#since they are not highly correlated.
data <- data %>%
  mutate(across(all_of(variables), ~ ifelse(. %in% codes, NA, .))) %>%
  na.omit()
```

```
#Exploratory Data Analysis
```

```
# Histogram of Age distribution by Age
```


```
ggplot(data, aes(x = Age, fill = Stroke)) +
  geom_histogram(position = "identity", alpha = 0.6, bins = 30) +
  labs(
    title = "Age Distribution by Stroke Status",
    fill = "Stroke"
  ) +
  theme_minimal(base_size = 14) +
  theme(
    legend.position = "top"
  )
```

Age Distribution by Stroke Status



```
# A density plot of sleeping hours by Stroke
ggplot(data, aes(x = `Hours Of Sleep`, fill = Stroke)) +
  geom_density(alpha = 0.6, color = NA) +
  coord_cartesian(xlim = c(0, 13)) +
  labs(
    title = "Distribution of the Sleep Hours by Stroke",
    fill = "Stroke"
  ) +
  theme_minimal(base_size = 12) +
  theme(
    legend.position = "top"
  )
```

Distribution of the Sleep Hours by Stroke



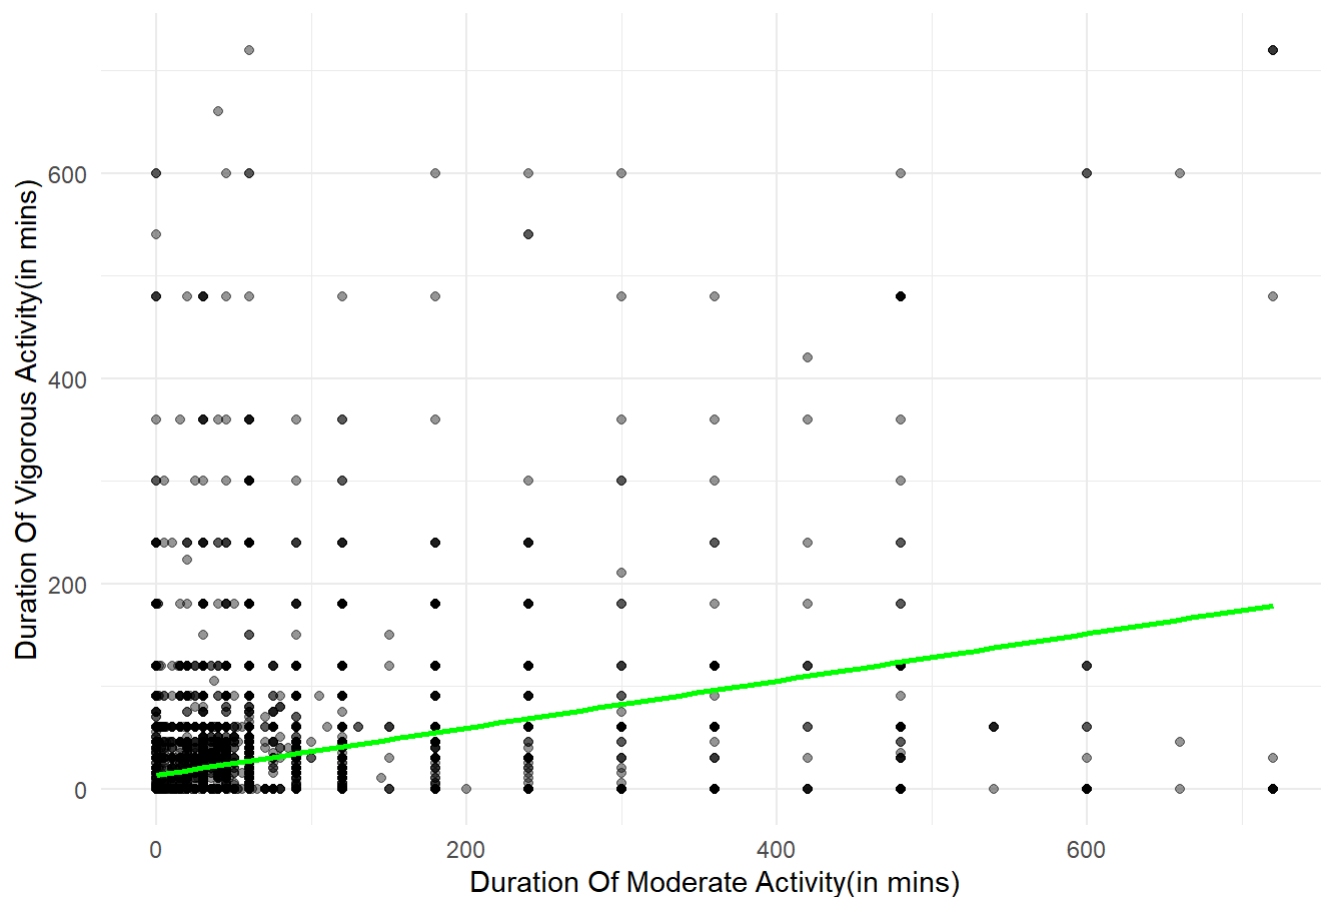
```
# Computing correlation between moderate & vigorous activity
correlation <- cor(
  data$`Duration Of Moderate Activity(in mins)`,
  data$`Duration Of Vigorous Activity(in mins)`
)
print(paste("Correlation (r) =", round(correlation, 5)))
```

```
## [1] "Correlation (r) = 0.31152"
```

```
# Plotting the above
ggplot(data, aes(
  x = `Duration Of Moderate Activity(in mins)`,
  y = `Duration Of Vigorous Activity(in mins)`
)) +
  geom_point(alpha = 0.4) +
  geom_smooth(method = "lm", se = FALSE,
             color = "green") +
  labs(
    title = "Moderate Activity vs Vigorous Activity (mins)",
  ) +
  theme_minimal()
```

```
## `geom_smooth()` using formula = 'y ~ x'
```


Moderate Activity vs Vigorous Activity (mins)



```
# Scaling the variables variables
data[variables] <- scale(data[variables])

# Splitting the data into train and test sets.
set.seed(42)
train_data <- createDataPartition(data$Stroke, p = 0.7, list = FALSE)
train_set <- data[train_data, ]
test_set <- data[-train_data, ]

# Increased 'Yes' class weight to 50
#to force the model to better recognize
#minority class during training.
# Chosen this setting of balance, to address the imbalance
# while avoiding extremes. This ratio improved
# stroke detection without excessive false alarms.
weights <- c("No" = 1, "Yes" = 50)
```

```
#PART 1: Linear SVM
set.seed(42)
tune_one <- tune(svm,
  Stroke ~ `Age` + Sex + `Hours Of Sleep` + `Hours Worked`
  + `Alcohol Consumption Days Per Year`
  + `Cigarettes Consumed Per Month`
  + `Duration Of Moderate Activity(in mins)`
  + `Duration Of Vigorous Activity(in mins)`,
  data      = train_set, kernel      = "linear",
  ranges    = list(cost = c(0.01, 0.1, 1)),
  class.weights = weights)
```

Using class weights to handle imbalance without dropping data

```
# Choosing the best linear svm from tuning
svm_one <- tune_one$best.model
svm_one
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Stroke ~ Age + Sex + `Hours Of Sleep` +
##   `Hours Worked` + `Alcohol Consumption Days Per Year` + `Cigarettes Consumed Per Month`
## +
##   `Duration Of Moderate Activity(in mins)` + `Duration Of Vigorous Activity(in mins)`,
##   data = train_set, ranges = list(cost = c(0.01, 0.1, 1)), kernel = "linear",
##   class.weights = weights)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1
##
## Number of Support Vectors:  8739
```

```
# Prediction and evaluation metrics
prediction_one <- predict(svm_one, test_set)
confusion_mat_one <- confusionMatrix(prediction_one, test_set$Stroke)
print(confusion_mat_one)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   No  Yes
##           No 3846  35
##           Yes 1650  89
##
##           Accuracy : 0.7002
##           95% CI : (0.688, 0.7121)
##           No Information Rate : 0.9779
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0567
##
##           McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.69978
##           Specificity : 0.71774
##           Pos Pred Value : 0.99098
##           Neg Pred Value : 0.05118
##           Prevalence : 0.97794
##           Detection Rate : 0.68434
##           Detection Prevalence : 0.69057
##           Balanced Accuracy : 0.70876
##
##           'Positive' Class : No
##
```

```
precision_one <- posPredValue(prediction_one, test_set$Stroke, positive = "Yes")
recall_one <- sensitivity(prediction_one, test_set$Stroke, positive = "Yes")
f1_one <- 2 * (precision_one * recall_one) / (precision_one + recall_one)
accuracy_one <- confusion_mat_one$overall["Accuracy"]
precision_one
```

```
## [1] 0.05117884
```

```
recall_one
```

```
## [1] 0.7177419
```

```
f1_one
```

```
## [1] 0.09554482
```

```
accuracy_one
```

```
## Accuracy
## 0.7001779
```


```

# Extracting variable importance from the fitted Linear SVM model
weight_vector <- as.numeric(t(svm_one$coefs) %*% svm_one$SV)
names(weight_vector) <- colnames(svm_one$SV)

# 'importance_df' consists of absolute weights and sort them in decreasing order
importance_df <- data.frame(
  Variable = names(weight_vector),
  Importance = abs(weight_vector)
)
importance_df <- importance_df[order(importance_df$Importance, decreasing = TRUE), ]

# Plotting the variable importance
ggplot(importance_df, aes(
  x = reorder(Variable, Importance),
  y = Importance
)) +
  geom_col(fill = "#619CFF") +
  coord_flip() +
  labs(
    title = "Variable Importance (Linear SVM)",
    x = NULL,
    y = "Importance"
  ) +
  theme_minimal(base_size = 13)

```



```
ggsave("variable_imp.png", width = 5, height = 5, dpi = 350, bg = "white")
```

Evaluated model with confusion matrix (Accuracy is 70%)

#PART 2: Radial SVM

```
tune_two <- tune(svm,
  Stroke ~ `Age` + Sex + `Hours Of Sleep` + `Hours Worked`
  + `Alcohol Consumption Days Per Year`
  + `Cigarettes Consumed Per Month`
  + `Duration Of Moderate Activity(in mins)`
  + `Duration Of Vigorous Activity(in mins)`,
  data          = train_set,
  kernel        = "radial",
  ranges        = list(cost = c(0.1, 1), gamma = c(0.01, 0.1)),
  class.weights = weights)
```

Choosing the best radial svm from tuning

```
svm_two <- tune_two$best.model
svm_two
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Stroke ~ Age + Sex + `Hours Of Sleep` +
##   `Hours Worked` + `Alcohol Consumption Days Per Year` + `Cigarettes Consumed Per Month`
##   +
##   `Duration Of Moderate Activity(in mins)` + `Duration Of Vigorous Activity(in mins)`,
##   data = train_set, ranges = list(cost = c(0.1, 1), gamma = c(0.01,
##     0.1)), kernel = "radial", class.weights = weights)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:  1
##
## Number of Support Vectors:  7352
```

Prediction and evaluation metrics

```
prediction_two <- predict(svm_two, test_set)
confusion_mat_two <- confusionMatrix(prediction_two, test_set$Stroke)
print(confusion_mat_two)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   No  Yes
##           No 3896  36
##           Yes 1600  88
##
##           Accuracy : 0.7089
##           95% CI : (0.6968, 0.7208)
##           No Information Rate : 0.9779
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0584
##
##           McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.70888
##           Specificity : 0.70968
##           Pos Pred Value : 0.99084
##           Neg Pred Value : 0.05213
##           Prevalence : 0.97794
##           Detection Rate : 0.69324
##           Detection Prevalence : 0.69964
##           Balanced Accuracy : 0.70928
##
##           'Positive' Class : No
##
```

```
precision_two <- posPredValue(prediction_two, test_set$Stroke, positive = "Yes")
recall_two<- sensitivity(prediction_two, test_set$Stroke, positive = "Yes")
f1_two<- 2 * (precision_two * recall_two) / (precision_two + recall_two)
accuracy_two<- confusion_mat_two$overall["Accuracy"]
precision_two
```

```
## [1] 0.0521327
```

```
recall_two
```

```
## [1] 0.7096774
```

```
f1_two
```

```
## [1] 0.09713024
```

```
accuracy_two
```

```
## Accuracy
## 0.7088968
```

Radial SVM achieved is equal to 70% accuracy on test set.

#PART 3: Polynomial SVM

```
tune_three <- tune(svm,
  Stroke ~ `Age` + Sex + `Hours Of Sleep` + `Hours Worked`
  + `Alcohol Consumption Days Per Year`
  + `Cigarettes Consumed Per Month`
  + `Duration Of Moderate Activity(in mins)`
  + `Duration Of Vigorous Activity(in mins)`,
  data          = train_set, kernel          = "polynomial",
  ranges        = list(cost = c(0.1,1), degree = c(3,4),
                        coef0 = c(0.5,1)),
  class.weights = weights)
```

Choosing the best polynomial svm from tuning

```
svm_three <- tune_three$best.model
svm_three
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Stroke ~ Age + Sex + `Hours Of Sleep` +
##   `Hours Worked` + `Alcohol Consumption Days Per Year` + `Cigarettes Consumed Per Month`
##   +
##   `Duration Of Moderate Activity(in mins)` + `Duration Of Vigorous Activity(in mins)`,
##   data = train_set, ranges = list(cost = c(0.1, 1), degree = c(3,
##     4), coef0 = c(0.5, 1)), kernel = "polynomial", class.weights = weights)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:    1
##    degree:   4
##   coef.0:   0.5
##
## Number of Support Vectors:  6850
```

Prediction and evaluation metrics

```
prediction_three<- predict(svm_three, test_set)
confusion_mat_three <- confusionMatrix(prediction_three, test_set$Stroke)
print(confusion_mat_three)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   No  Yes
##           No 3996  36
##           Yes 1500  88
##
##           Accuracy : 0.7267
##           95% CI : (0.7148, 0.7383)
##           No Information Rate : 0.9779
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0645
##
##           McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.72707
##           Specificity : 0.70968
##           Pos Pred Value : 0.99107
##           Neg Pred Value : 0.05542
##           Prevalence : 0.97794
##           Detection Rate : 0.71103
##           Detection Prevalence : 0.71744
##           Balanced Accuracy : 0.71838
##
##           'Positive' Class : No
##
```

```
precision_three<- posPredValue(prediction_three, test_set$Stroke, positive = "Yes")
recall_three<- sensitivity(prediction_three, test_set$Stroke, positive = "Yes")
f1_three <- 2 * (precision_three * recall_three) / (precision_three + recall_three)
accuracy_three <- confusion_mat_three$overall["Accuracy"]
precision_three
```

```
## [1] 0.05541562
```

```
recall_three
```

```
## [1] 0.7096774
```

```
f1_three
```

```
## [1] 0.1028037
```

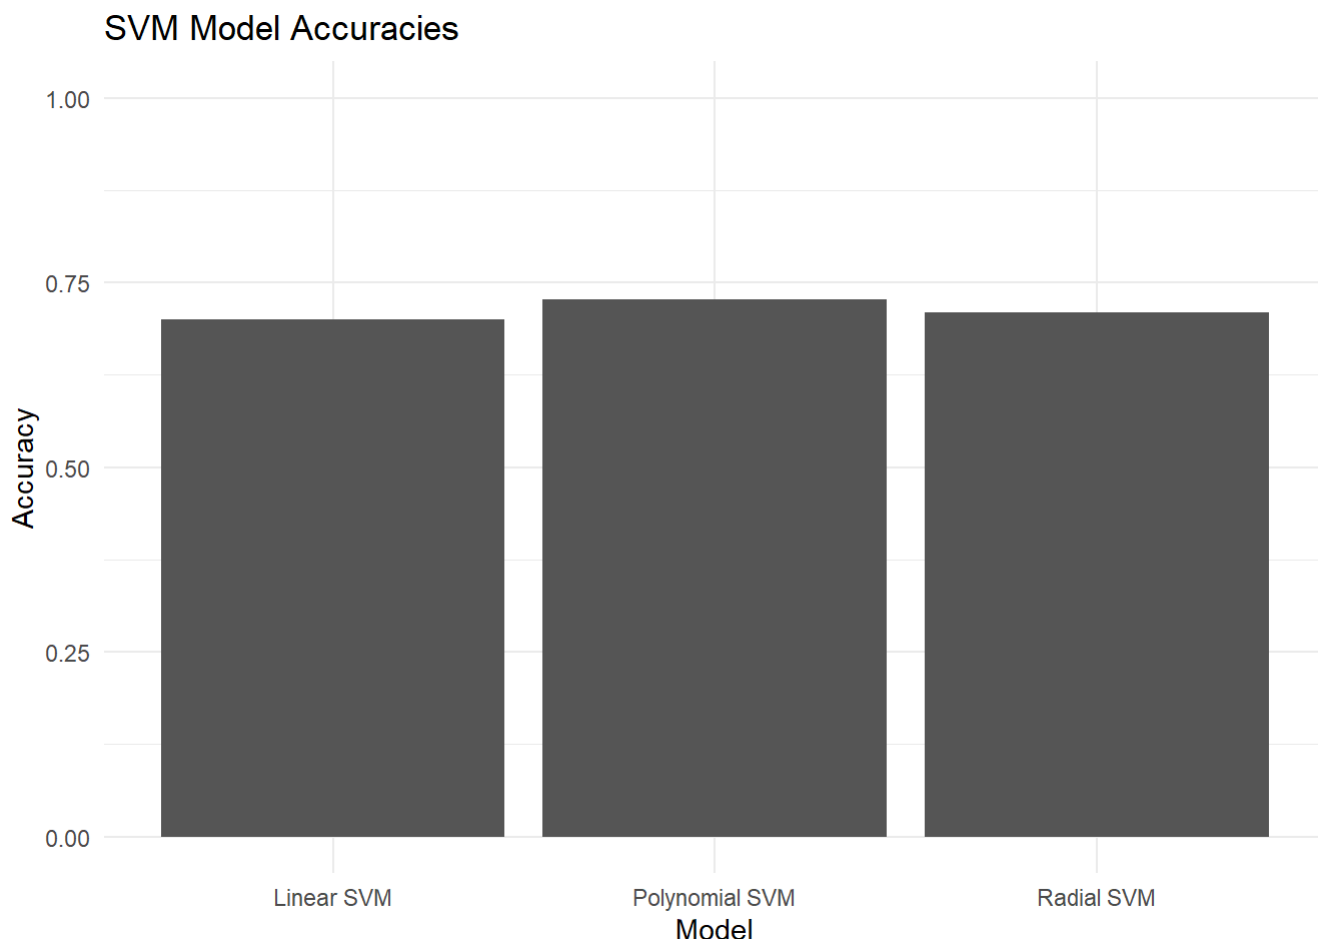
```
accuracy_three
```

```
## Accuracy
## 0.7266904
```

Polynomial SVM achieved is close to 73% accuracy on test set.


```
# Comparing the accuracy results by plotting
model_results <- data.frame(
  Model = c("Linear SVM", "Radial SVM", "Polynomial SVM"),
  Accuracy = c(accuracy_one, accuracy_two, accuracy_three)
)

ggplot(model_results, aes(x = Model, y = Accuracy)) +
  geom_col() +
  ylim(0, 1) +
  labs(title = "SVM Model Accuracies", x = "Model", y = "Accuracy") +
  theme_minimal()
```



```
# creating the model evaluation table
evaluation_table <- data.frame(
  Model      = c("Linear SVM", "Radial SVM", "Polynomial SVM"),
  Accuracy   = c(accuracy_one, accuracy_two, accuracy_three),
  Precision  = c(precision_one, precision_two, precision_three),
  Recall     = c(recall_one, recall_two, recall_three),
  F1_Score   = c(f1_one, f1_two, f1_three)
)

# rounding off the numeric values to present
evaluation_table <- evaluation_table %>%
  mutate(across(where(is.numeric), ~ round(., 3)))

print(evaluation_table)
```

##	Model	Accuracy	Precision	Recall	F1_Score
## 1	Linear SVM	0.700	0.051	0.718	0.096
## 2	Radial SVM	0.709	0.052	0.710	0.097
## 3	Polynomial SVM	0.727	0.055	0.710	0.103

From the above plot of SVM accuracies comparison, polynomial model has performed the highest, about 73 % , radial performs next best , and then linear svm performs good, does purely linear separation


```
# Subset training data with the two strong predictors
train_subset <- train_set %>%
  select(`Hours Worked`, Age, Stroke)
```

```
# Fitting the Linear SVM with the two predictors and one response variables.
svm_four <- svm(Stroke ~ `Hours Worked` + Age,
               data = train_subset,
               kernel = "linear",
               cost = 1,
               scale = TRUE,
               class.weights = weights)
```

```
# Creating a grid and then plotting
x_seq <- seq(min(train_subset$`Hours Worked`), max(train_subset$`Hours Worked`), length.out = 200)
y_seq <- seq(min(train_subset$Age), max(train_subset$Age), length.out = 200)
grid <- expand.grid(`Hours Worked` = x_seq, Age = y_seq)
grid$Prediction <- predict(svm_four, grid)

ggplot() +
  geom_tile(data = grid, aes(x = `Hours Worked`, y = Age, fill = Prediction), alpha = 0.3) +
  geom_point(data = train_subset, aes(x = `Hours Worked`, y = Age, shape = Stroke, color = Stroke), size = 2.0) +
  scale_fill_manual(values = c("No" = "#FFC0CB", "Yes" = "#90EE90")) +
  scale_color_manual(values = c("No" = "deeppink3", "Yes" = "forestgreen")) +
  labs(title = "Linear SVM Decision Boundary",
       x = "Hours Worked",
       y = "Age",
       fill = "SVM Region",
       color = "Stroke Status",
       shape = "Stroke Status") +
  theme_minimal(base_size = 16) +
  theme(
    plot.title = element_text(hjust = 0.5, face = "bold"),
    legend.position = "right"
  )
```

Linear SVM Decision Boundary



```
ggsave("linear_svm2.png", width = 4, height = 4, dpi = 350, bg = "white")
```

```
#From the above plots, polynomial model has performed the highest,  
#about 73 % , radial performs next best ,  
#and then linear svm performs good, does purely linear separation.
```

```
#RESULTS: The linear SVM, with best cost 1 , has 8739 support vectors,  
#suggesting high complexity, and it has an accuracy of 70%.  
#It is poor at detecting strokes, precision is 5.1%;
```

```
#The radial SVM, with best cost 1 and gamma 0.1 ,  
#has 7352 support vectors,so its efficient,  
#and it has an accuracy of 70.89%.  
#It is slightly better than linear svm,  
#but still poor at detecting strokes, precision is 5.21%;
```

```
#The polynomial SVM is the best among all,  
#with best cost 1, degree 4, and coef0 0.5 , has 6850 support vectors,  
#and it has an accuracy of 72.66%.  
#It is still suffers from low precision like the rest of the two models,  
#for "Yes" , but has better F1-score.
```

```
#The minority class is "yes" stroke and  
#all the three models perform poorly in this one.
```